

## FINDING EXACT AND APPROXIMATE BLOCK STRUCTURES FOR ILU PRECONDITIONING\*

YOUSEF SAAD†

**Abstract.** Sparse matrices which arise in many applications often possess a block structure that can be exploited in iterative and direct solution methods. These block-matrices have as their entries small dense blocks with constant or variable dimensions. Block versions of incomplete LU factorizations which have been developed to take advantage of such structures give rise to a class of preconditioners that are among the most effective available. This paper presents general techniques for automatically determining block structures in sparse matrices. A standard “graph compression” algorithm used in direct sparse matrix methods is considered along with two other algorithms which are also capable of unraveling approximate block structures.

**Key words.** graph compression, incomplete LU factorization, block ILU preconditioners, hash-based graph compression, cosine-based graph compression, indistinguishable nodes

**AMS subject classifications.** 65F10, 65N06

**PII.** S1064827501393393

**1. Introduction.** ILU-type preconditioners combined with Krylov subspace accelerators are currently among the most effective iterative techniques for solving large, sparse, irregularly structured, linear systems of equations. Incomplete LU factorizations, which consist of performing a Gaussian elimination and dropping some fill-in, give rise to an important class of efficient and relatively robust preconditioners. Among these, the subclass of “block ILU” techniques is known by practitioners to be extremely powerful for some types of problems, though this is seldom emphasized in the standard numerical linear algebra literature. A block ILU method is one that treats (small) dense submatrices of the matrix  $A$  as single entities. The classical case is when the coefficient matrix  $A$  is a discretization of a partial differential equation in which a fixed number of variables, say  $l$ , is associated with each mesh point. In this case, the matrix can be viewed as a “block-matrix,” i.e., a matrix whose entries are dense  $l \times l$  submatrices. This situation is fairly common, for example, in computational fluid dynamics, where it is typical for  $l$  to be equal to 4 in two dimensions, and 5 in three dimensions. In two dimensions, the four variables may represent density ( $\rho$ ), scaled energy ( $\rho E$ ), and two scaled velocity components of the fluid ( $\rho u, \rho v$ ). In three dimensions, an additional velocity component is added.

Block ILU preconditioning applied to such matrices gives rise to a fast and fairly robust iterative solution procedure; see, for example, [6]. It is known that it is always a good strategy to use a block version of ILU instead of a scalar version when this is possible. An obvious advantage of such techniques is the savings in storage, since most column indices and pointers for the block entries are avoided. Indeed, the usual block sparse row format used in SPARSKIT [11] employs a compressed sparse row (CSR) format for the resulting  $(n/l) \times (n/l)$  block matrix, but each value in the CSR

---

\*Received by the editors August 8, 2001; accepted for publication (in revised form) June 6, 2002; published electronically January 31, 2003. This work was supported in part by the Army Research Office under grant DAAD19-00-1-0485, and in part by the NSF under grants NSF-INT 0003274 and NSF/ACI-0000443.

<http://www.siam.org/journals/sisc/24-4/39339.html>

†Department of Computer Science and Engineering, University of Minnesota, 4-192 EE/CS Building, 200 Union Street S.E., Minneapolis, MN 55455 (saad@cs.umn.edu).

format becomes an array of size  $l \times l$ . A more important advantage is the potential gain obtained from computing with dense blocks and using level 3 basic linear algebra subroutine (BLAS3)-based computations.

It is easy to generalize the constant block-size format to a variable block-size format; see, for example, the SPARSKIT package [11]. Such generalizations are important because many applications lead to matrices with variable blocks.

The problem of finding a block structure for a matrix has been considered from a number of different angles in the past. For example, O'Neil and Szyld [10] proposed a scheme named PABLO, which consists of reordering the matrix in such a way as to obtain diagonal blocks that are dense. During sparse direct solution methods, the eliminating column tends to propagate its pattern, leading to sets of columns referred to as "indistinguishable nodes." A technique for identifying these sets was proposed by Ashcraft [1]. It is clear that the sets of indistinguishable nodes provide a simple way of blocking a matrix. Such techniques are widespread in packages for reordering sparse matrices in sparse direct solution packages; see, e.g., the nested dissection ordering in Metis (see [8]) and the MUMPS package [3], among others.

What motivated the topic of this paper is that many of the matrices that have a relatively large number of nonzero elements per row have a "near block structure." This means that if we allow the nonzero pattern to expand a little, by making a few zero entries part of the nonzero pattern, we easily obtain a variable block matrix. Finding an "approximate" block structure is not as easy as finding an exact blocking. We propose two algorithms and show that the process of grouping the rows/columns is in general much less expensive than that of performing the standard block ILU factorization.

**2. Graph compression and matrix blocking.** In this section, we consider only matrices with symmetric patterns. For illustration, consider the matrix represented below, where an  $x$  represents a nonzero element:

$$(1) \quad A = \left[ \begin{array}{cc|cc|ccc|c} x & x & 0 & 0 & x & x & x & 0 \\ x & x & 0 & 0 & x & x & x & 0 \\ \hline 0 & 0 & x & x & 0 & 0 & 0 & x \\ 0 & 0 & x & x & 0 & 0 & 0 & x \\ \hline x & x & 0 & 0 & x & x & x & 0 \\ x & x & 0 & 0 & x & x & x & 0 \\ x & x & 0 & 0 & x & x & x & 0 \\ \hline 0 & 0 & x & x & 0 & 0 & 0 & x \end{array} \right].$$

The above matrix has a clear variable block structure, with blocks of size 2, 2, 3, and 1, respectively. The four sets of nodes into which  $V$  is partitioned are  $Y_1 = \{1, 2\}$ ,  $Y_2 = \{3, 4\}$ ,  $Y_3 = \{5, 6, 7\}$ , and  $Y_4 = \{8\}$ . To avoid confusion, we took an example in which each block consists of contiguous columns, but it is clear that this is not required. For example, another natural partition consists of the two subsets  $Z_1 = Y_1 \cup Y_3$  and  $Z_2 = Y_2 \cup Y_4$ .

Let us call  $\mathcal{P}$  the partition into the four sets  $Y_i$ ,  $i = 1, \dots, 4$ . Since the matrix has a symmetric pattern, the adjacency graph  $G = (V, E)$  of the above matrix can be more succinctly represented by its quotient graph [7] denoted by  $G/\mathcal{P} = \{V_{\mathcal{P}}, E_{\mathcal{P}}\}$  and defined by

$$V_{\mathcal{P}} = \{Y_1, \dots, Y_p\}, \quad E_{\mathcal{P}} = \{(Y_i, Y_j) \mid \exists v \in Y_i, w \in Y_j \text{ s.t. } (v, w) \in V\}.$$

In the above example, the adjacency matrix for the quotient graph is the  $4 \times 4$  matrix

$$\begin{bmatrix} x & 0 & x & 0 \\ 0 & x & 0 & x \\ x & 0 & x & 0 \\ 0 & x & 0 & x \end{bmatrix}.$$

One can consider that the entry in position  $(i, j)$  is a dense block of dimension  $|Y_i| \times |Y_j|$ , where  $|X|$  is the cardinality of the set  $X$ . Automatically finding the partition  $\mathcal{P}$  is useful in many different ways in sparse matrix computations. In sparse direct solution methods, two vertices are said to be indistinguishable if they have the same pattern. Finding a blocking of the matrix in variable block format is clearly equivalent to grouping its vertices into subsets of indistinguishable nodes.

**2.1. Hash-based algorithms.** Blockings of the form shown in the previous subsection are relatively inexpensive to unravel. A technique used in several existing packages associates a key, or “checksum” value, to each row. The simplest such key used in [1] is the following:

$$(2) \quad \text{key}(u) = \sum_{(u,w) \in E} w.$$

If the checksum of two rows is different, then clearly the rows have a different pattern. If they are the same, then they may or may not have the same pattern, and thus an explicit comparison of the row patterns becomes necessary. Sorting the keys in a first step facilitates the process. At any given step of the block construction, a row will be compared with all rows succeeding it (in the order of the keys). For as long as the keys are the same, a comparison of the patterns is made with the pattern of  $\text{row}_i$ . The algorithm is described in some detail below. In the description,  $K(i)$  is the checksum key shown above for each row  $i$ , and  $\text{Group}(i)$  is a representative array:  $\text{Group}(i) = k$  means that  $i$  belongs to group  $k$ , unless  $k$  is equal to  $-1$ , in which case  $i$  is the representative of the group.

ALGORITHM 2.1. HASH-BASED COMPRESSION.

1. Compute the keys  $K(u)$  for each vertex  $u$ .
2. Set  $\text{Group}(i) = -1$  for  $i = 1, \dots, n$ .
3. List in  $Krow(1 : n)$  the rows of  $A$  in increasing order of the keys.
4. For  $i = 1 : n$  Do
5.      $row = Krow(i)$
6.     For  $j = i + 1 : n$  Do
7.          $rowj = Krow(j)$ ;
8.         If  $\text{key}(rowj) \neq \text{key}(row)$ , break;
9.         If  $\text{Group}(rowj) == -1$ ,
10.             If  $\text{pattern}(rowj) == \text{pattern}(row)$ ,  $\text{Group}(rowj) = row$ ;
11.         EndIf.
12.     EndDo.
13. EndDo.

In order to give an idea of the performance of the above algorithm, we gathered statistics on a few runs for a sample of 14 matrices with different degrees of blocking and various sizes. Some generic information about the matrices is shown in Table 1. All matrices are available from the Matrix-Market<sup>1</sup> except for PENALTY, which is

<sup>1</sup><http://math.nist.gov/MatrixMarket/>

TABLE 1  
Information on the 14 matrices used for the current tests.

Matrix	Type	$n$	$nnz$
BCSSTK10	RSA	1086	11578
BCSSTK11	RSA	1473	17857
BCSSTK12	RSA	1473	17857
BCSSTK16	RSA	4884	147631
RAEFSKY1	RUA	3242	294276
RAEFSKY2	RUA	3242	294276
RAEFSKY3	RUA	21200	1488768
RAEFSKY4	RUA	19779	1328611
BARTHT1A	RUA	14075	481125
BARTHT2A	RUA	14075	1311725
FIDAP006	RUA	1651	49479
FIDAP023	RUA	1409	43481
FIDAP028	RUA	2603	77653
PENALTY	RUA	35186	625828

TABLE 2  
Performance of the hash-based algorithm on a test sample of 14 sparse matrices.

Matrix	Compression statistics			VBILUK Time	
	V. Cmpr	E. Cmpr	Time	levf = 0	levf = 2
BCSSTK10	2.30	5.52	0.00	0.04	0.05
BCSSTK11	1.89	3.46	0.01	0.07	0.17
BCSSTK12	1.89	3.46	0.01	0.06	0.17
BCSSTK16	2.73	7.45	0.08	0.49	2.83
RAEFSKY1	3.31	12.74	0.08	0.43	4.49
RAEFSKY2	3.31	12.74	0.14	0.47	4.61
RAEFSKY3	8.00	64.00	0.43	1.95	4.25
RAEFSKY4	2.93	8.75	0.38	2.07	25.28
BARTHT1A	5.00	25.00	0.16	0.63	1.15
BARTHT2A	5.01	25.07	0.35	1.71	10.56
FIDAP006	1.35	2.08	0.01	0.18	0.88
FIDAP023	1.20	1.47	0.01	0.21	1.20
FIDAP028	1.45	2.55	0.02	0.21	1.09
PENALTY	2.00	4.00	0.28	1.01	2.10

obtained from the Diffpack package [9, 5]. Note that the BCSSTK matrices are all symmetric, whereas the solvers being used do not take advantage of symmetry. The other matrices are all nonsymmetric. In the table,  $n$  is the dimension of the matrix, and  $nnz$  represents the total number of nonzero elements. The types “RSA” and “RUA” correspond to the labels “real symmetric assembled” and “real unsymmetric assembled” used in the Harwell–Boeing collection. For RSA matrices, only the lower triangular part is stored, and  $nnz$  represents its number of nonzero elements.

Table 2 shows some statistics for the hash-based algorithm. The second and third columns of the table are the vertex and edge compression rates, respectively, achieved by the algorithm. These are the ratios  $|V|/|V_P|$  and  $|E|/|E_P|$ . The next column shows the time required to execute the algorithm. The last two columns show, for reference, the time it takes to compute the variable block level-of-fill ILU with fill-levels of 0 and 2, respectively. As can be seen, the compression times are quite low—in fact, in many cases negligible—relative to the time for the VBILU(k) factorization.

Before looking at alternative methods, we discuss potential improvements to the checksum-based technique. Clearly, the idea of checksums is similar to that of hash functions, and the term hash was explicitly used in [2]. A row is “hashed” into a given

value. Two rows hashing into the same value is like a “collision” in the hash table, and this makes it necessary to check the patterns. Similarly to hashing, it is possible to improve the performance of the method by clever choices of the hash function. An ideal hash function would assign a different value for each different row pattern that occurs. If we knew that the hash function had this property, then there would be no need to compare patterns. Assuming that the vertices are labeled from 1 to  $n$ , a one-to-one hash function is

$$(3) \quad \text{hash}(u) = \sum_{(u,w) \in E} 2^{w-1},$$

which is simply the integer represented by the pattern of the row viewed as a binary number (0 for a zero element, 1 for a nonzero element). The above “perfect” hash function is not practical because it leads to huge numbers that may not even be machine-representable. If nothing is exploited about the pattern, then this is the only function that is one-to-one from the  $2^n$  possible patterns to the set  $1, \dots, 2^n$ . If we wish to reduce the comparisons to only those rows that have the same pattern, then clearly we can exploit what is known from constructing hash functions to reduce the number of collisions to a minimum. However, since for most matrices the time required to find indistinguishable nodes is negligible relative to the time it takes to solve the system, this is an issue that has not received much attention. This can be easily understood from the results of Table 2. Therefore, potential improvements to the simple choice (2) may exist but will not be explored in this paper.

**2.2. The cosine algorithm.** Consider now a slightly modified pattern obtained from the sample matrix  $A$  in (1). It can, for example, happen that the matrix in (1) is “filtered” of its small elements and loses three entries in the process, resulting in a matrix that is no longer block:

$$(4) \quad A = \left[ \begin{array}{cc|cc|ccc|c} x & x & 0 & 0 & x & x & x & 0 \\ x & x & 0 & 0 & x & x & 0 & 0 \\ \hline 0 & 0 & x & x & 0 & 0 & 0 & x \\ 0 & 0 & x & x & 0 & 0 & 0 & x \\ \hline x & x & 0 & 0 & x & 0 & x & 0 \\ x & x & 0 & 0 & x & x & x & 0 \\ x & 0 & 0 & 0 & x & x & x & 0 \\ \hline 0 & 0 & x & x & 0 & 0 & 0 & x \end{array} \right].$$

The checksum algorithm would be able to find only one nontrivial block, namely (2,3), all others being of size 1. However, in this case, it is clear that we still would like to consider the matrix as block—with the same blocking as in (1). By doing so, we make a small sacrifice in memory usage since a few zero entries are considered nonzeros to pad rows into desirable patterns. We refer to these “zero nonzero” entries as fill-ins, for lack of a better term. Ideally, we would need a method that can find an “approximate block structure.” It is not easy to extend the hash-based algorithm to handle this case; that would require a hash function that preserves the proximity of patterns, in the sense that close patterns will result in close hash values. In addition, when rows which have the same key (or possibly nearby keys) must be compared, we will need to count the number of matching column locations.

As an alternative to hash-based algorithms, we consider a technique which compares angles of rows (or columns). It may at first appear expensive to compare all




FIG. 1. Computation of the inner product of  $c_i$  with all columns of  $C^T$ .

the angles of the rows of a matrix with each other. However, a good implementation is key to making the method effective.

Let us call  $C$  the adjacency matrix related to  $A$ . This is a matrix which has the same pattern as  $A$  and whose nonzero values are equal to one. The main idea is to compute the  $i$ th row of  $CC^T$  or, to be more accurate, the upper triangular part of this row. Entry  $(i, j)$  in this row is the inner product of row  $i$  with row  $j$ . We need consider only those entries  $(i, j)$  with  $j > i$ . The inner product will give the cosine between row  $i$  and row  $j$ , and if the corresponding angle is small enough,  $j$  will be added to the “group”  $i$ . This is indicated by setting  $Group(j) = i$ . The group of node  $i$  itself is set to  $-1$  to indicate that this entry is a representative of a group. This is repeated for  $i = 1, \dots, n$ , but when a row is already assigned ( $Group(i) > 0$ ), the row is skipped.

Key to this algorithm is the computation of  $c_i C^T$ , which is done rowwise in the usual way that one proceeds when multiplying two sparse matrices in a rowwise structure, with a few important variations. This process is illustrated in Figure 1. Note that we describe the algorithm for general matrices with possibly nonsymmetric patterns. The goal is to find groups of rows which have similar patterns. When seeking a block structure for a matrix, it is often assumed that the pattern is symmetric, although for block ILU, a symmetric pattern is not required. We will come back to this issue in a separate section.

The basic cost of computing all inner products of row  $i$  with all columns  $C^T$  (i.e., rows of  $C$ ) is the sum of the number of nonzero entries of each row involved in this inner product:

$$\sum_{j=1}^{|c_i|} |c_j^T|.$$

Here  $c_j^T$  denotes the  $j$ th row of  $C^T$ , and  $|\cdot|$  is the cardinality. This is basically the cost of computing the  $i$ th row of the upper triangular part of  $CC^T$ , except that there are no floating point operations made. However, there are several simplifications and improvements to this basic scheme which will substantially reduce the cost. The first is that row  $i$  is skipped if  $i$  has already been assigned to an existing group. As the algorithm progresses, many rows may already be assigned, and thus this could lead to substantial savings. In addition, when computing the inner products, note that if column  $j$  with  $j > i$  has already been assigned, it can be ignored in the calculation, since its inner product will not be of any use.

A final improvement comes from the fact that only entries  $j > i$  of  $c_i C^T$  need to be computed. This can be easily exploited, provided that the entries of each row of  $C^T$  are sorted by increasing column numbers. When  $C^T$  is actually obtained by transposing  $C$  with standard sparse transposition algorithms such as the one in SPARSKIT [11], for example, this property is true. Assume that this is the case. Then when “adding” row  $k$  to the current working row of  $c_i C^T$ , we will proceed backward from the rightmost column and move down until a column entry that is  $\leq i$  is encountered, in which case the loop is stopped. The final algorithm is sketched below. The notation  $nz_A(i)$  means the number of nonzero entries in the  $i$ th row of  $A$ . The meaning of  $Group(i)$  is identical with that seen in Algorithm 2.1.

ALGORITHM 2.2. COSINE-BASED COMPRESSION.

Input: pattern matrix  $C$  and tolerance  $\tau$ . Output: set data structure for blocks.

1. Compute the pattern matrix  $C^T$ .
2. Set  $Group(i) = -1$  and  $Count(i) = 0$  for  $i = 1, \dots, n$ .
3. For  $i = 1 : n$  and if  $Group(i) == -1$ , Do
4.     For  $\{j \mid c_{ij} \neq 0\}$  Do
5.         Let  $row = j$ th row of  $C^T$ .
6.         For  $k = nz_{C^T}(j)$  downto 1 Do
7.             Let  $col = row(k)$ .
8.             if  $col \leq i$ , break;
9.             if  $(Group(col) == -1)$ ,  $Count(col) = Count(col) + 1$ ;
10.         EndDo.
11.     EndDo.
12.     For each  $col$  such that  $Count(col) \neq 0$  Do
13.         If  $Count(col)^2 > \tau^2 * nz_A(i) * nz_A(col)$ , then
14.              $Group(col) = i$ ;
15.         EndIf.
16.          $Count(col) = 0$ .
17.     EndDo.
18. EndDo.

The scalar  $\tau$  in line 13 is a threshold parameter which determines acceptance of a new row into a group. A new row is added to a group when the cosine of the angle it makes with the reference row is larger than  $\tau$ . So  $\tau$  is typically a number that is close to one.

Table 3 shows statistics that are similar to those shown in Table 2 for the hash-based algorithm. A grouping tolerance of  $\tau = 0.8$  was used for the tests. An additional column labeled “compression efficiency” is added, which represents the ratio of the original number of nonzero elements over the number of nonzero elements of the block matrix (taking fill-ins into account). An ideal compression that introduces no fill-in has an efficiency of one, which would correspond to a percentage value of 100% in Table 3. There are two observations worth noting. First, several matrices see a substantial increase in their edge compression rate, by introducing a moderate fill-in of less than 10%. For example, edge compression for BCSTK11 goes from 3.46 to 8.67 with an efficiency of 90.39%. A more substantial gain is seen for FIDAP023, whose edge compression rate goes from 1.47 to 6.11 with a similar efficiency. A second observation is that the cost of the algorithm increases substantially from that of the hash-based technique. In the worst case (BARTH2A), the time increases nearly fivefold. However, in most cases it still remains inferior to the cost of the least expensive block LU factorization, i.e., block ILU(0). Compared with the cost of the more realistic ILU factorization such as VBILU(2), the cost is essentially still negligible. This is illustrated in Table 3.

TABLE 3

Performance of the cosine-based algorithm on a test sample of 14 sparse matrices.

Matrix	Compression statistics				VBILUK Time	
	V. Cmpr	E. Cmpr	Cmpr eff	Time	levf = 0	levf = 2
BCSSTK10	3.58	9.87	80.50%	0.02	0.04	0.06
BCSSTK11	3.02	8.67	90.39%	0.03	0.05	0.10
BCSSTK12	3.02	8.67	90.39%	0.03	0.06	0.11
BCSSTK16	4.31	15.56	79.24%	0.36	0.54	2.29
RAEFSKY1	3.64	14.11	93.61%	0.53	0.45	4.41
RAEFSKY2	3.64	14.11	93.61%	0.53	0.45	4.40
RAEFSKY3	8.64	69.16	95.23%	1.28	2.10	4.40
RAEFSKY4	4.31	13.85	79.49%	1.81	2.50	24.89
BARTHT1A	5.01	25.07	99.87%	0.41	0.63	1.14
BARTHT2A	6.20	33.52	88.35%	1.71	2.01	10.78
FIDAP006	3.02	6.70	93.37%	0.06	0.09	0.26
FIDAP023	3.02	6.11	90.11%	0.05	0.08	0.26
FIDAP028	3.18	7.10	87.81%	0.09	0.15	0.40
PENALTY	2.07	4.15	97.67%	0.71	0.99	2.06

**2.3. A hybrid algorithm.** As is shown by the experiments seen so far, the angle algorithm is more expensive than hash-based techniques. It is easy to combine both algorithms into one, resulting in a cost that is closer to that of hash-based methods for cases in which a good blocking already exists. Since the quotient graph will be used, the technique introduced in this section is restricted to matrices with symmetric patterns. Simply put, the idea is to do a first pass with the hash-based algorithm to detect any “exact” block structure and then to perform the angle algorithm on the quotient graph. The cost of this second pass, which works on the typically smaller quotient graph, is likely to be a small addition to the cost of the initial blocking by hashing. In the second pass, the algorithm scans each nonassigned row again to determine whether it can be added to an existing group.

ALGORITHM 2.3. HYBRID METHOD FOR COMPRESSION.

Input: pattern matrix  $C$  and tolerance  $\tau$ . Output: set data structure for blocks.

1. Call Algorithm 2.1 to find an initial group  $Group_0$  assignment. Set  $Group = Group_0$ .
2. Obtain the adjacency matrix  $C$  of the quotient graph  $G_{\mathcal{P}}$ , using the original labeling.
3. For  $i = 1 : n$  and if  $Group(i) == -1$ , Do
4.     For  $\{j \mid c_{ij} \neq 0\}$  Do
5.         Let  $row = j$ th row of  $C^T$  and  $s = |Group_0(j)|$ .
6.         For  $k = nz_{C^T}(j)$  downto 1 Do
7.             Let  $col = row(k)$ .
8.             If  $col \leq i$ , break;
9.             If  $(Group_0(col) == -1)$ ,  $Count(col) = Count(col) + s$ ;
10.         EndDo.
11.     EndDo.
12.     For each  $col$  such that  $Count(col) \neq 0$  Do
13.         If  $Count(col)^2 > \tau^2 * nz_A(i) * nz_A(col)$ , then
14.              $Group(col) == i$ ;
15.         EndIf.
16.          $Count(col) = 0$ .
17.     EndDo.
18. EndDo.



TABLE 4  
Performance of the hybrid algorithm on a test sample of 14 sparse matrices.

Matrix	Compression statistics				VBILUK Time	
	V. Cmpr	E. Cmpr	Cmpr eff	Time	levf = 0	levf = 2
BCSSTK10	3.58	9.87	80.50%	0.01	0.05	0.06
BCSSTK11	3.02	8.67	90.39%	0.02	0.06	0.11
BCSSTK12	3.02	8.67	90.39%	0.03	0.05	0.11
BCSSTK16	4.31	15.56	79.24%	0.17	0.55	2.30
RAEFSKY1	3.64	14.11	93.61%	0.15	0.45	4.36
RAEFSKY2	3.64	14.11	93.61%	0.16	0.45	4.36
RAEFSKY3	8.63	69.16	95.23%	0.60	2.11	4.39
RAEFSKY4	4.31	13.85	79.49%	0.77	2.52	24.90
BARTHT1A	5.01	25.07	99.87%	0.26	0.63	1.14
BARTHT2A	6.20	33.52	88.35%	0.58	2.00	10.69
FIDAP006	3.02	6.70	93.37%	0.05	0.09	0.26
FIDAP023	3.02	6.11	90.11%	0.06	0.08	0.26
FIDAP028	3.18	7.10	87.81%	0.07	0.14	0.41
PENALTY	2.07	4.15	97.67%	0.57	1.00	2.08

The second part of the algorithm is similar to that of Algorithm 2.2. One notable difference is that the increment in line 9 is now  $s$ , the size of the  $j$ th group obtained in the first path. The above algorithm works on the quotient graph but weighs the nonzero entries differently so that we will get the same inner products as in Algorithm 2.2.

Table 4 shows statistics that are similar to those shown in Tables 2 and 3. The same grouping tolerance of  $\tau = 0.8$  as in Table 3 was used. Observe that, as expected, the compression statistics are identical with those of the cosine-based algorithm. This is achieved at a cost that is, in most cases, much closer to the cost of the hash-based algorithm. Another observation is that, as expected, the costs of the cosine and hybrid algorithms are close when the hash-based method does not achieve a good compression. This is the case for FIDAP023, for example.

**2.4. Fill-in analysis.** An important question regarding the angle-based compression algorithm is related to its “memory efficiency.” Blocking is useful in reducing overall execution times, by effectively exploiting BLAS3 computations and possibly by yielding better ILU techniques. On the other hand, when inexact blocking is performed, we may have to introduce many nonzero elements to pad rows in order to obtain dense blocks. Consider, for example, the two rows

$$\begin{array}{cccccccc} 1 & 0 & 0 & 1 & 1 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 1 & 0 & 0 & 1, \end{array}$$

where it is assumed that the first row (denoted by  $r$ ) is the reference row in the cosine-based algorithm and we denote by  $u$  the second row. Denote by  $\{v\}$  the set associated with the row  $v$ , which comprises all column indices in the row, and by  $|v|$  the cardinality of this set. The squared 2-norms of  $r$  and  $u$  are their numbers of nonzero elements,  $|r| = 6$  and  $|u| = 4$ , respectively. Their inner product  $\langle r, u \rangle$ , which is equal to 4 in this case, is the number of matching 1’s in the two patterns. If it is decided that  $u$  will be part of the group lead by  $r$ , according to the algorithm, then the new pattern of  $r$ , the union of the patterns of  $u$  and  $r$ , becomes

$$1 \ 1 \ 0 \ 1 \ 1 \ 1 \ 0 \ 1.$$

As was already emphasized, the algorithm does not update these patterns, as to do so is expensive. Instead, it uses the same initial pattern to determine the angles between

rows, the rationale being that the difference in the patterns will likely be small. It is useful to note, however, that the union of all the row patterns in the group will be the final pattern of the group. The size of the union of  $\{r\}$  and  $\{u\}$  is easily seen to be given by  $|r_{new}| = |r| + |u| - \langle r, u \rangle$ , the sum of the sizes of  $\{r\}$  and  $\{u\}$  minus the size of the intersection  $\{r\} \cap \{u\}$ . Assume that the algorithm ends up adding rows  $u_k$ ,  $k = 1, \dots, s_r - 1$ , to the group  $\{r\}$ . After each of these additions, the size of the new group pattern is given by

$$|r_k| = |r_{k-1}| + |u_k| - \langle u_k, r_{k-1} \rangle,$$

starting with  $r_0 \equiv r$ . Since the pattern of  $r_{k-1}$  includes the pattern of  $r$ , we have

$$\langle u_k, r_{k-1} \rangle \geq \langle u_k, r \rangle \geq \tau \sqrt{|u_k| |r|}$$

and therefore

$$|r_k| \leq |r_{k-1}| + |u_k| - \tau \sqrt{|u_k| |r|},$$

which leads to the following upper bound for the size of the final row  $\tilde{r} \equiv r_{s_r}$  obtained from the union with the patterns  $u_1, \dots, u_{s_r-1}$ :

$$(5) \quad |\tilde{r}| \leq |r| + \sum_{k=1}^{s_r-1} \left( |u_k| - \tau \sqrt{|u_k| |r|} \right) = |r| + |r| \sum_{k=1}^{s_r-1} \left( \frac{|u_k|}{|r|} - \tau \sqrt{\frac{|u_k|}{|r|}} \right).$$

To simplify the analysis, we now assume that all of the rows  $u_k$  that are added to the group have the same length, which is also the length of  $r$ . In this case, the ratio  $|u_k|/|r|$  is equal to one, and (5) simplifies into

$$(6) \quad |\tilde{r}| \leq |r| [1 + (s_r - 1)(1 - \tau)].$$

This gives an upper bound for the expansion of the reference row  $r$  under the simplifying assumption that all rows in the group have the same number of nonzero elements.

The integer  $s_r - 1$  represents the number of rows that have been added to the group, and thus  $s_r$  is just the size of the group. Consider the situation in which all the rows  $u_k$  are modified so that they have the same pattern as  $\tilde{r}$ . This corresponds to forming the “block” associated with group  $r$  by including all fill-ins. The block can be viewed as a dense matrix of dimension  $s_r \times s_r$ , which is reordered into the large matrix. It is clear that the number of rows in this dense block does not exceed the number of nonzero entries in the final row  $\tilde{r}$ , since this row includes the nonzero elements of the first row of the block. This is illustrated in the following example, in which rows 2 and 5 are added to the first group to form the final group:

$$\begin{array}{cccccc} 1 & 1 & 0 & 0 & 1 & 0 & 1 \\ 1 & 1 & 0 & 0 & 1 & 0 & 1 \\ \vdots & \vdots & & & & \vdots & \\ 1 & 1 & 0 & 1 & 1 & 0 & 1 \end{array} \longrightarrow \begin{array}{cccccc} 1 & 1 & 0 & 1 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 1 & 0 & 1 \\ \vdots & \vdots & & & & \vdots & \\ 1 & 1 & 0 & 1 & 1 & 0 & 1 \end{array}.$$

In the example, the final (filled) block consists of a dense  $3 \times 3$  matrix which is obtained from rows and columns (1, 2, 5). The number of rows in this block, which is 3, does not exceed the number of nonzeros in the first filled row of the (sparse global) matrix,

which is 5, because this contains at least the entries  $a_{11}, a_{12}, a_{15}$  of the dense block. Therefore, we can state that

$$(7) \quad s_r \leq |\tilde{r}|,$$

and this leads to

$$|\tilde{r}| \leq |r| [1 + (|\tilde{r}| - 1)(1 - \tau)].$$

If we now assume that  $1 - \tau$  is small enough that  $1 - |r|(1 - \tau) < 1$ , then the above inequality yields

$$(8) \quad |\tilde{r}| \leq |r| \frac{1}{1 - |r|(1 - \tau)}.$$

The above upper bound will not be sharp in general because of its reliance on (7) to bound the size of each group. In a typical situation, (7) provides a rough estimate. However, the expansion due to fill-in can be estimated from the inequality (6), provided we again make the assumption that each row  $u$  in the group has the same number of nonzero elements as the original reference row  $r$  (of the original matrix) and that, in addition, each group has a size that does not exceed a certain  $\beta$ . In this case, denoting by  $\tilde{A}$  the matrix  $A$  after blocking is applied to it (and fill-in is incorporated), and by  $g_1, g_2, \dots, g_m$  all the groups found by the algorithm, we can write

$$\text{nnz}(\tilde{A}) = \sum_{g_r} |g_r| |\tilde{r}| \leq (1 + (\beta - 1)(\tau - 1)) \sum_{g_r} |g_r| |r| = (1 + (\beta - 1)(\tau - 1)) \text{nnz}(A).$$

This can be restated as follows.

**PROPOSITION 2.1.** *Assume that during Algorithm 2.2 (and 2.3) each group is such that the number of nonzero elements in its member rows is constant, and let  $\beta$  be the maximum block-size found by these algorithms. If  $\tilde{A}$  is the blocked matrix which includes fill-ins due to blocking, then*

$$(9) \quad 1 \leq \frac{\text{nnz}(\tilde{A})}{\text{nnz}(A)} \leq 1 + (\beta - 1)(\tau - 1).$$

Notice in particular that, as expected, there is no fill-in when the blocks are all of size  $\beta = 1$ , i.e., when no blocking is discovered.

Although the maximum block-size found by the algorithm is not known, it is intuitively clear that it cannot be arbitrarily large. In fact, from (7) and (8) we can see that for  $\tau$  sufficiently small we have

$$\beta \leq |r_{\max}| \frac{1}{1 - |r_{\max}|(1 - \tau)},$$

in which  $|r_{\max}|$  is the maximum number of nonzero elements per row in  $A$ . For the same reasons that (8) is not a sharp inequality, the above upper bound is also unlikely to be sharp.

**3. Matrices with nonsymmetric patterns.** Compression can also be used to unravel exploitable block structure for matrices with nonsymmetric patterns, though

quotient graphs can no longer be used. For example, a column (only) block structure such as the one in the matrix

$$A = \left[ \begin{array}{cc|cc|ccc|c} x & x & 0 & 0 & x & x & x & 0 \\ x & x & x & x & x & x & x & 0 \\ 0 & 0 & x & x & 0 & 0 & 0 & x \\ x & x & 0 & 0 & 0 & 0 & 0 & x \\ x & x & x & x & x & x & x & 0 \\ 0 & 0 & x & x & x & x & x & 0 \\ x & x & 0 & 0 & x & x & x & 0 \\ x & x & 0 & 0 & 0 & 0 & 0 & x \end{array} \right]$$

could clearly be exploited in a row-oriented block ILU.

There are, however, a few difficulties. Whereas in the symmetric case a diagonal block is either a zero block or a square dense block, this is not guaranteed in the nonsymmetric case, since the blocks are not necessarily square. As an example, the second block column in the above matrix has as a diagonal block a (singular)  $2 \times 2$  matrix with a zero second row. This situation does not arise if all the diagonal elements of  $A$  are nonzero and if exact blocking is performed ( $\tau = 1$  for the angle and the hybrid algorithms). One way to avoid this difficulty is obviously to symmetrize the pattern, and this will be a good strategy in the rather common case in which the pattern is nearly symmetric. If a small number of diagonal entries are zero, another strategy is to always treat diagonal entries as nonzero.

For matrices whose patterns are very far from being symmetric, compression is best done in two stages. The column-compressed graph for the above matrix consists of considering each set of indistinguishable columns as a node. This yields the following  $8 \times 4$  adjacency matrix:

$$\left[ \begin{array}{cccc} x & 0 & x & 0 \\ x & x & x & 0 \\ 0 & x & 0 & x \\ x & 0 & 0 & x \\ x & x & x & 0 \\ 0 & x & x & 0 \\ x & 0 & x & 0 \\ x & 0 & 0 & x \end{array} \right].$$

Now a row-based compression can be applied to the above matrix, yielding the row-partition

$$Y_1 = \{1, 7\}, \quad Y_2 = \{2, 5\}, \quad Y_3 = \{3\}, \quad Y_4 = \{8\}.$$

The row and column blockings may be quite different. In fact, the number of groups may not even be the same as the number of groups found for the columns. However, now each block of the blocked matrix is a dense block in the case in which exact blocking is used. Computations can be reorganized in a complete or incomplete LU factorization to take advantage of this nonsymmetric blocking.

**4. Numerical tests with VBILUK.** The goal of the numerical experiments in this section is to illustrate how the methods described earlier can be integrated within a variable block ILU preconditioner with fill level (VBILUK). We have implemented and tested only a level-of-fill block ILU technique, with variable blocks. All codes have

been written in C, and, unless otherwise stated, the experiments have been conducted on a PC with an Intel Pentium II processor with a clock speed of 450MHz and 500MB of main memory. The timings in the following experiments were obtained with a  $-O3$  optimization directive during compilation (in contrast with those of Tables 2, 3, and 4 of the previous sections, where no optimization was used).

The principle of the block ILU preconditioner is straightforward, and so details will be omitted. It suffices to say that the ILU factorization is conceptually performed on the quotient matrix, and the levels-of-fill are computed for the corresponding quotient graph. The inversion of the diagonal blocks is performed via a truncated SVD factorization for better stability. Although this implementation does not seem to have been discussed in the literature, the simpler version of block ILU( $k$ ) with constant block size is well established; see, for example, [6]. One problem with using a block ILU factorization with variable blocks is precisely the fact that such a blocking is not known in advance or is difficult to handle practically. With automatic blocking, we need to provide only one parameter, namely, the tolerance for grouping. If a hash-based algorithm is used, no parameter is required.

In this test we consider the impact of the grouping parameter  $\tau$  on the performance of VBILUK. For this we took the same sample of the matrices tested in the previous tables. However, we removed two matrices and added a new one. We removed PENALTY and RAEFSKY4 because the methods did not converge for these two hard problems. We added VENKAT25, a matrix of dimension  $n = 62,424$  with  $nnz = 1,717,792$  nonzeros, which originates from an unstructured two-dimensional Euler solver.<sup>2</sup> VENKAT25 has a natural blocking with a block size of 4. All systems are solved by forming an artificial right-hand side and taking a random initial guess. The iteration was stopped when the residual norm was decreased by 10 orders of magnitude or when an iteration count of 300 was reached. The accelerator GMRES(60) was used in all tests.

Recall that as  $\tau$  decreases from 1, the blocks are likely to become larger. A rule of thumb is that we may expect the number of iterations to improve slightly. This is observed in general but, as Table 5 shows, not always. As can be seen, there is at least one case (FIDAP028) for which the hash-based algorithm (which is equivalent to setting  $\tau = 1$ ) yields convergence whereas the other cases do not. For all other cases, the grouping obtained by the hybrid algorithm yields better or comparable iteration counts for VBILU(0) and VBILU(2). Thus, from this small sample of experiments it is observed that using larger blocks by blocking near-matching rows does help convergence. A clear advantage of larger blocks is obviously a better utilization of computational hardware. In sparse direct solvers the strategy of using dense computations at the expense of storage is already well known; see, e.g., [3]. Table 6 shows the execution times associated with the results of Table 5. As can be seen, the fact that ILU uses possibly much larger blocks does not penalize performance. In fact, in many cases, the execution is at least marginally better than that obtained from the pure blocking yielded by the hash-based algorithm.

In an attempt to understand the convergence failures encountered in the previous experiments, we show in Table 7 a simple indicator of the quality of the preconditioner, namely, the norm  $\|L^{-1}AU^{-1} - I\|_1$  for those smaller matrices for which problems were encountered. In some of the cases shown, this is a fairly good indicator of the difficulty when taken in a relative way, i.e., when comparing preconditioners obtained for the same matrix. For example, for FIDAP028, VBILU(0) has serious difficulties

<sup>2</sup>Also available from the matrix market at <http://math.nist.gov/MatrixMarket/>

TABLE 5

Iteration counts for GMRES(60) preconditioned with VBILU( $k$ ), for  $k = 0, 2$ , on a collection of linear systems.

Matrix	Hash	$\tau = .70$	$\tau = .75$	$\tau = .80$	$\tau = .90$	$\tau = .95$
VBILU(0)						
BCSSTK10	119	46	40	37	54	54
BCSSTK11	300	80	86	80	85	140
BCSSTK12	300	80	86	80	85	140
BCSSTK16	47	42	39	40	47	44
RAEFSKY1	36	51	46	36	36	36
RAEFSKY2	45	56	49	45	45	45
RAEFSKY3	78	62	62	62	78	78
BARTHT1A	88	69	73	88	88	88
BARTHT2A	52	50	51	52	51	52
FIDAP006	300	300	300	300	300	300
FIDAP023	98	177	95	58	57	97
FIDAP028	106	300	300	300	300	127
VENKAT25	238	239	238	238	238	238
VBILU(2)						
BCSSTK10	14	14	13	9	14	14
BCSSTK11	30	26	30	30	30	30
BCSSTK12	30	26	30	30	30	30
BCSSTK16	16	13	14	14	16	16
RAEFSKY1	19	21	21	19	19	19
RAEFSKY2	22	24	24	22	22	22
RAEFSKY3	50	49	49	49	50	50
BARTHT1A	30	30	30	30	30	30
BARTHT2A	23	19	21	22	23	23
FIDAP006	36	101	300	113	37	36
FIDAP023	25	22	17	17	18	19
FIDAP028	300	30	46	113	300	300
VENKAT25	87	87	87	87	87	87

in all cases except for  $\tau = 1$  and  $\tau = 0.95$ , and in these cases, the error norm is much smaller than for most of the other cases. However, this is not uniform. For example, it is very surprising that VBILU(2) yields convergence for FIDAP006 when the preconditioning error norm is  $2.2e+19$ , which is incidentally the highest norm achieved for this matrix.

An interesting comment is in regards to matrices which already have a natural blocking. The times obtained for solving some of these systems hardly vary, because the number of GMRES iterations remains the same for all compressions. Two examples are VENKAT25 and BARTH2A, using  $k = 2$ . This is because the compressions obtained here are close to optimal, in that they are hard to improve by additional groupings.

To give an idea of the potential gains that can be achieved from a block version over a point version of the same preconditioner, we compare in Table 8 the preprocessing and iteration times for VBILUK(2) and point ILUK(2) on two matrices under the same test conditions as those in the previous tests. The block version used the blocking provided by the hash-based algorithm (equivalent to setting  $\tau = 1.0$  in the other two algorithms). As is known, the point and block versions of ILU( $k$ ) are mathematically equivalent. As a result, the number of iterations should be identical for both, as the column "Its" in the table confirms. The column "Pr-Mem." reports the number of memory locations required by the preconditioners, which as expected are also identical. The iterations times ("Pr-sec.") are close for both methods. In

TABLE 6

Iteration times for solving a collection of linear systems by GMRES(60) preconditioned with VBILU( $k$ ), for  $k = 0, 2$ .

Matrix	Hash	$\tau = .70$	$\tau = .75$	$\tau = .80$	$\tau = .90$	$\tau = .95$
VBILU(0)						
BCSSTK10	1.04	0.34	0.28	0.26	0.45	0.45
BCSSTK11	4.91	0.93	1.02	0.93	1.06	1.87
BCSSTK12	4.95	0.93	1.04	1.00	1.03	1.88
BCSSTK16	3.81	2.97	2.76	2.83	3.61	3.39
RAEFSKY1	2.31	3.21	3.07	2.27	2.29	2.40
RAEFSKY2	2.95	3.62	3.24	3.09	3.02	2.96
RAEFSKY3	22.76	18.81	18.61	19.38	22.60	23.02
BARTHT1A	10.72	8.66	9.12	10.90	10.93	10.72
BARTHT2A	13.29	14.77	13.88	13.71	13.26	13.42
FIDAP006	8.66	4.96	5.17	5.35	5.67	6.14
FIDAP023	2.95	2.78	1.51	0.92	0.97	1.79
FIDAP028	4.31	8.05	8.00	8.43	8.84	4.35
VENKAT25	157.59	155.76	151.38	152.91	157.96	153.56
VBILU(2)						
BCSSTK10	0.13	0.11	0.10	0.07	0.13	0.13
BCSSTK11	0.68	0.41	0.47	0.46	0.50	0.54
BCSSTK12	0.68	0.42	0.47	0.46	0.49	0.54
BCSSTK16	2.43	1.48	1.66	1.71	2.24	2.30
RAEFSKY1	2.88	2.85	3.00	3.11	2.87	2.95
RAEFSKY2	3.34	3.29	3.41	3.33	3.37	3.35
RAEFSKY3	19.30	18.96	18.77	19.17	19.04	19.29
BARTHT1A	4.55	4.50	4.62	4.55	4.68	4.54
BARTHT2A	10.74	10.35	10.22	10.06	10.73	10.81
FIDAP006	2.14	4.46	7.97	8.59	3.44	1.25
FIDAP023	1.61	0.52	0.41	0.41	0.48	0.62
FIDAP028	25.55	1.18	7.21	13.26	14.63	16.61
VENKAT25	67.33	65.82	65.41	65.89	66.26	65.84

TABLE 7

Norms of  $\|L^{-1}AU^{-1} - I\|_1$  relative to preconditioners in Tables 5 and 6 for a few matrices.

Matrix	Hash	$\tau = .70$	$\tau = .75$	$\tau = .80$	$\tau = .90$	$\tau = .95$
VBILU(0)						
BCSSTK10	5.677e+03	1.450e+04	1.530e+01	1.652e+02	1.453e+03	1.453e+03
BCSSTK11	1.710e+05	5.331e+02	4.738e+02	4.946e+02	5.233e+02	3.509e+04
BCSSTK12	1.710e+05	5.331e+02	4.738e+02	4.946e+02	5.233e+02	3.509e+04
FIDAP006	4.468e+03	7.150e+03	2.019e+15	2.662e+16	2.019e+15	9.805e+04
FIDAP023	5.482e+06	6.842e+11	5.376e+11	5.376e+11	5.140e+06	2.880e+06
FIDAP028	9.808e+02	1.108e+03	7.589e+15	1.258e+16	5.645e+16	1.728e+04
VBILU(2)						
BCSSTK10	2.179e+00	4.077e+00	4.387e+00	1.606e+00	2.182e+00	2.182e+00
BCSSTK11	3.966e+04	5.351e+04	1.188e+05	3.132e+04	3.787e+04	3.961e+04
BCSSTK12	3.966e+04	5.351e+04	1.188e+05	3.132e+04	3.787e+04	3.961e+04
FIDAP006	2.061e+05	2.220e+19	4.491e+18	2.741e+16	2.019e+15	5.825e+04
FIDAP023	1.508e+07	4.447e+11	5.376e+11	5.376e+11	3.126e+06	2.880e+06
FIDAP028	1.207e+18	4.212e+02	7.589e+15	1.425e+16	6.029e+17	1.191e+18

the current implementation of the code, matrix-vector operations do not take advantage of blocking in the iteration phase. LAPACK [4] routines are extensively used in the construction phase and in the forward-backward sweeps. Thus, it is likely that additional gains in time can be made for the iteration phase by better code optimization. However, the gains made when constructing the preconditioner are substantial,

TABLE 8  
Performance of a block and point ILU( $k$ ) on two matrices.

Matrix	VBILUK(2)				ILUK(2)			
	Pr-Mem.	Its	Its sec.	Pr-sec.	Pr-Mem.	Its	Its sec.	Pr-sec.
BARTHT1A	958775	30	4.53	1.02	958775	30	4.44	3.01
RAEFSKY3	2810240	50	19.33	3.41	2810240	50	21.57	15.84

reaching a reduction by a factor of about 4.6 for RAEFSKY3.

**5. Conclusion.** The methods presented in this paper have as their primary goal to automate blocking in block ILU factorizations. In order to allow for imperfect blocking, i.e., for blocking in which the entries in the blocks are allowed to be zeros, it is useful to be able to group rows according to the nearness of their patterns. This can be achieved by using the angle between the rows as a measure of nearness. The algorithms to accomplish this are inexpensive when compared, for example, to the cost of BILU(0), the least expensive factorization.

There are other possible applications of the blocking techniques presented here. In preconditioning methods, they can be combined with algebraic recursive multilevel solvers (ARMS) [12]. In a multilevel ILU context, the successive approximate Schur complements that are generated can become quite dense, and blocking can have a good performance pay-off. In fact, this is precisely the strategy utilized in sparse direct solution software to obtain flop rates that are far superior to those of iterative solvers.

**Acknowledgements.** The numerical experiments were conducted by Na Li, a graduate student in computer science at the University of Minnesota. The block ILU programs used in this paper were translated into C from FORTRAN-77 codes written by Andrew Chapman (see [6]). The author benefited from discussions on the topic of this paper with Patrick Amestoy and John Gilbert during a visit to CERFACS, France, supported by an NSF/INRIA grant. The Minnesota Supercomputer Institute provided computing resources and an excellent environment for conducting this research.

#### REFERENCES

- [1] C. ASHCRAFT, *Compressed graphs and the minimum degree algorithm*, SIAM J. Sci. Comput., 16 (1995), pp. 1404–1411.
- [2] P. R. AMESTOY, T. A. DAVIS, AND I. S. DUFF, *An approximate minimum degree ordering algorithm*, SIAM J. Matrix Anal. Appl., 17 (1996), pp. 886–905.
- [3] P. R. AMESTOY, I. S. DUFF, AND J.-Y. L'EXCELLENT, *MUMPS Multifrontal Massively Parallel Solver*, version 2.0, Technical report TR/PA/98/02, CERFACS, Toulouse, France, 1998.
- [4] E. ANDERSON, Z. BAI, C. BISCHOF, J. DEMMEL, J. DONGARRA, J. DUCROZ, A. GREENBAUM, S. HAMMARLING, A. MCKENNEY, S. OSTROUCHOV, AND D. SORENSSEN, *LAPACK Users' Guide*, SIAM, Philadelphia, PA, 1992.
- [5] X. CAI AND A. ODEGÅRD, *Parallel simulation of 3d nonlinear acoustic fields on a Linux-cluster*, in Proceedings of the IEEE International Conference on Cluster Computing (CLUSTER 2000), Saxony, Germany, 2000; also available at <http://www.ifi.uio.no/~xingca/xc-papers.shtml>.
- [6] A. CHAPMAN, Y. SAAD, AND L. WIGTON, *High-order ILU preconditioners for CFD problems*, Int. J. Numer. Methods Fluids, 33 (2000), pp. 767–788.
- [7] J. A. GEORGE AND J. W. LIU, *Computer Solution of Large Sparse Positive Definite Systems*, Prentice-Hall, Englewood Cliffs, NJ, 1981.
- [8] G. KARYPIS AND V. KUMAR, *A fast and high quality multilevel scheme for partitioning irregular graphs*, SIAM J. Sci. Comput., 20 (1998), pp. 359–392.



- [9] Å. ODEGÅRD, P. FOX, S. HOLM, AND A. TVEITO, *Finite element modelling of pulsed Bessel beams and x-waves using Diffpack*, in Proceedings of 25th International Imaging Symposium, Bristol, United Kingdom, 2000, pp. 59–64.
- [10] J. O'NEIL AND D. B. SZYLD, *A block ordering method for sparse matrices*, SIAM J. Sci. Comput., 11 (1990), pp. 811–823.
- [11] Y. SAAD, *SPARSKIT: A Basic Tool Kit for Sparse Matrix Computations*, Technical report RIACS-90-20, Research Institute for Advanced Computer Science, NASA Ames Research Center, Moffet Field, CA, 1990.
- [12] Y. SAAD AND B. SUCHOMEL, *ARMS: An Algebraic Recursive Multilevel Solver for General Sparse Linear Systems*, Numer. Linear Algebra Appl., 9 (2002), pp. 359–378.